

# Vahana.jl Workshop

Part I - Vahana.jl Basics & First Model Implementation

Steffen Fürst

Scaling Complexity Workshop  
Berlin, 03.12.2024

# Introducing Vahana.jl

- ▶ Vahana is an open source HPC framework designed for **large-scale ABMs**, particularly those focusing on complex social systems.
- ▶ Vahana has a focus on **network-based models**, although it is allowed to include a spatial dimension.
- ▶ Vahana simplifies the development phase with its **interactive capabilities**, using Julia's REPL.
- ▶ Once developed, a Vahana simulation can be **run parallelized** without additional actions from the developer.
- ▶ Models must be formulated as a **Graph Dynamic System** (GDS), which imposes certain constraints.

# (Synchronous) Graph Dynamical Systems (GDS)

- ▶ A GDS involve a **finite graph** with a state assigned to each vertex and a **vertex function** that update the state of a vertex based on the states of its neighboring vertices.
- ▶ In a synchronous GDS (SyGDS), all vertex state updates happen **simultaneously**.
- ▶ A SyGDS can be interpreted as a **generalized cellular automata**, where the neighborhood is determined by the network instead of the position of the cell.
- ▶ When an ABM is implemented as a GDS, the **vertices represent the individual agents** of the model.

# Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph**  $G_t = (V_t, E_t)$ .
- ▶ Vertices of **different types**  $\delta_v \in \Delta_V$ , and a mapping function  $\pi_v : V_t \rightarrow \Delta_V$ .
- ▶ Likewise edges of **different types**  $\delta_e \in \Delta_E$ , and a mapping function  $\pi_e : E_t \rightarrow \Delta_E$ .

# Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph**  $G_t = (V_t, E_t)$ .
- ▶ Vertices of **different types**  $\delta_v \in \Delta_V$ , and a mapping function  $\pi_v : V_t \rightarrow \Delta_V$ .
- ▶ Likewise edges of **different types**  $\delta_e \in \Delta_E$ , and a mapping function  $\pi_e : E_t \rightarrow \Delta_E$ .
- ▶ For each  $\delta_v$  and  $\delta_e$  a different **state space**  $\Theta_\delta = X_{\delta,1} \times \dots \times X_{\delta,n_\delta}$ , where each  $X_{\delta,i}$  is itself a set (like e.g.  $\mathbb{N}$ ).
- ▶ At least one **transition (vertex) function**  $f_{\delta_v,i} : G'_{v,t} \rightarrow (V'_{v,t+1}, E'_{v,t+1})$ , where  $G'_{v,t}$  is the 1-neighborhood of vertex  $v \in G_t$ .
- ▶ The source  $v_s$  and target  $v_t$  of an  $e \in E'_{v,t+1}$  must be vertices of  $G'_{v,t} \cup V'_{v,t+1}$ .

# Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph**  $G_t = (V_t, E_t)$ .
- ▶ Vertices of **different types**  $\delta_v \in \Delta_V$ , and a mapping function  $\pi_v : V_t \rightarrow \Delta_V$ .
- ▶ Likewise edges of **different types**  $\delta_e \in \Delta_E$ , and a mapping function  $\pi_e : E_t \rightarrow \Delta_E$ .
- ▶ For each  $\delta_v$  and  $\delta_e$  a different **state space**  $\Theta_\delta = X_{\delta,1} \times \dots \times X_{\delta,n_\delta}$ , where each  $X_{\delta,i}$  is itself a set (like e.g.  $\mathbb{N}$ ).
- ▶ At least one **transition (vertex) function**  $f_{\delta_v,i} : G'_{v,t} \rightarrow (V'_{v,t+1}, E'_{v,t+1})$ , where  $G'_{v,t}$  is the 1-neighborhood of vertex  $v \in G_t$ .
- ▶ The source  $v_s$  and target  $v_t$  of an  $e \in E'_{v,t+1}$  must be vertices of  $G'_{v,t} \cup V'_{v,t+1}$  **or a raster vertex**.

# Constructing $G_{t+1}$

To construct  $G_{t+1}$ , we take the following actions:

- ▶ For each vertex  $v \in V_t$  we **apply**  $f_{\pi_v(v), i}$ .
- ▶ **Combine** the results:  $G_{t+1} = (\bigcup_{v \in V_t} V'_{v, t+1}, \bigcup_{v \in V_t} E'_{v, t+1})$ .
- ▶ **Remove** all  $e \in E_{t+1}$  with a source  $v_s$  or target  $v_t$  that is not in  $V_{t+1}$ .

# Constructing $G_{t+1}$

To construct  $G_{t+1}$ , we take the following actions:

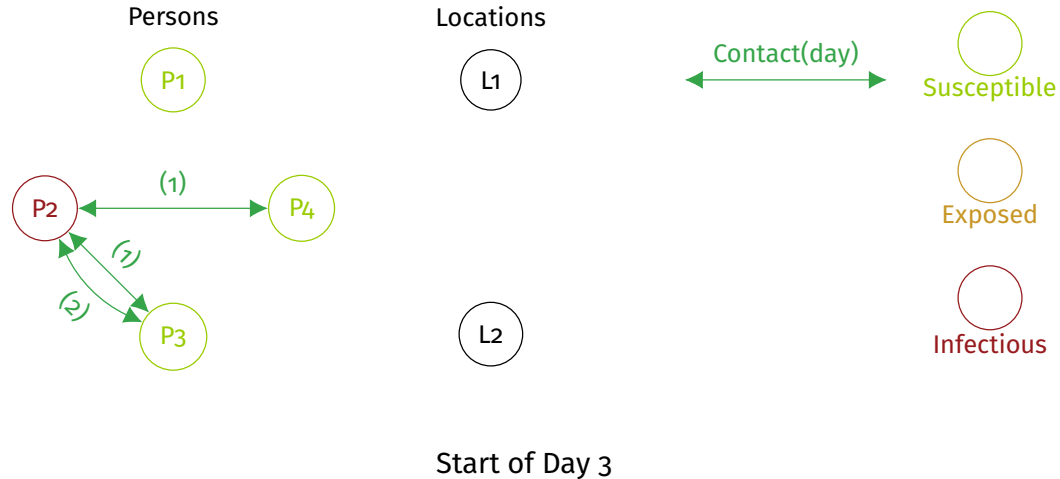
- ▶ For each vertex  $v \in V_t$  we **apply**  $f_{\pi_v(v),i}$ .
- ▶ **Combine** the results:  $G_{t+1} = (\bigcup_{v \in V_t} V'_{v,t+1}, \bigcup_{v \in V_t} E'_{v,t+1})$ .
- ▶ **Remove** all  $e \in E_{t+1}$  with a source  $v_s$  or target  $v_t$  that is not in  $V_{t+1}$ .

It would be inefficient and cumbersome to require that even **invariant parts of the graph** must be reconstructed, so:


- ▶ Users specify which vertex/edge types are **mutable** during transitions.
- ▶ There exist an option to retain existing agents/edges of a type, **only adding new ones**.




## Example: A epidemiological model



# Example: A epidemiological model




# Example: A epidemiological model



Locations: Infect persons & add contacts

# Example: A epidemiological model



Persons: Update health status

# REPL example

```
julia> show_agent(sim, Person; neighborstate = true)
Id / Local Nr: 0x01000000000085e35 / 548405
State:
  health=susceptible
  age=41
  quarantine=Quarantine(false, -1, -1)
Edge(s):
  MaybeEvent
    from:                edge.state:
    0x0200000000007691a MaybeEvent(0, ACTSTART, HOME)
    0x0200000000007691a MaybeEvent(32520, ACTEND, HOME)
    ... (16 not shown)
  Home
    from:
    0x0200000000007691a
  Knows
    from:                neighborstate:
    0x04000000000000001 capacity=0, traced=0, reported_infections=0,
```

# Parallel Simulations

- ▶ Vahana.jl uses the **Message Passing Interface** (MPI) for communication across processes.
- ▶ But it **shields users** from complexities of MPI.
- ▶ At the end of the initialization phase, the actual graph is **automatically partitioned** and distributed to the processes.
- ▶ Each **process only knows a subset** of the graph representing simulation state.
- ▶ Vahana.jl uses **MPI-3 shared memory** to avoid copies of agent state on same node.

# MATSim Episim


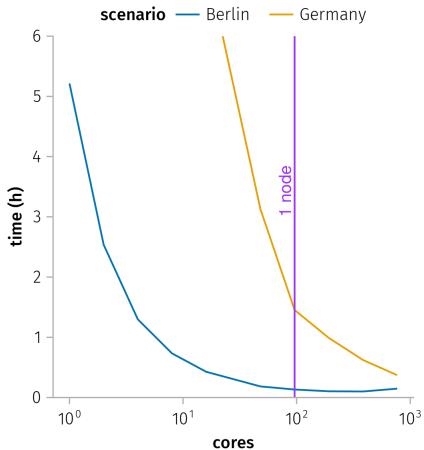
- ▶ MATSim Episim is an **epidemiological model** designed to predict the spread of infectious diseases within a population.
- ▶ It simulates individual agents moving between various **locations** such as their homes, workplaces, or schools.
- ▶ During the COVID-19 pandemic, this model was regularly used to **advise the German federal government** by providing reports containing evaluations of the **impact of various interventions**.
- ▶ Each report required **thousands of simulations** based on mobile-phone data for Berlin and Cologne, considering the activities of 25% of a city's population.

# MATSim Episim Scenarios

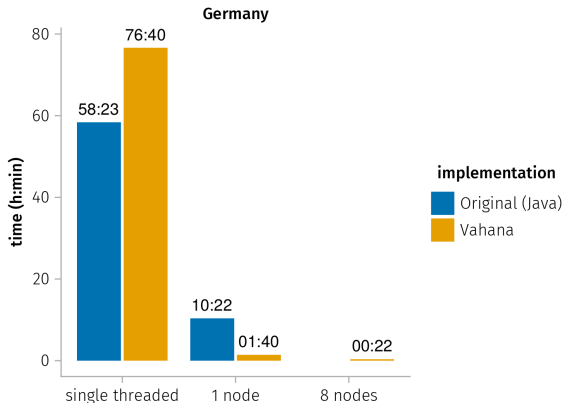
	Berlin	Germany
Persons	1,013,973	16,334,073
Locations	889,958	11,538,167
Location Visits	6,814,215	91,534,255



# Parallelization Performance of Vahana's Episim Implementation



# Parallelization Performance Compared to Original Implementation



With Vahana, simulations for the whole of Germany becomes practical.

# Hegselmann-Krause (HK) Opinion Model

- ▶ Finite number  $n$  of agents
- ▶ Time is **discrete**:  $t = 0, 1, \dots$
- ▶ Real number  $x_i(t)$  represents the **opinion** of agent  $i$
- ▶ There is a **confidence bound**  $\epsilon > 0$ , opinions with a difference greater than  $\epsilon$  are ignored.
- ▶ Opinions are updated **synchronously** according to

$$x_i(t+1) = \frac{1}{|\mathcal{N}_i(t)|} \sum_{j \in \mathcal{N}_i(t)} x_j(t)$$

$$\text{where } \mathcal{N}_i(t) = \{j : \|x_j(t) - x_i(t)\| \leq \epsilon\}$$