

Vahana.jl Workshop

Part I - Vahana.jl Basics & First Model Implementation

Steffen Fürst

Scaling Complexity Workshop
Berlin, 03.12.2024

Introducing Vahana.jl

- ▶ Vahana is an open source HPC framework designed for **large-scale ABMs**, particularly those focusing on complex social systems.
- ▶ Vahana has a focus on **network-based models**, although it is allowed to include a spatial dimension.
- ▶ Vahana simplifies the development phase with its **interactive capabilities**, using Julia's REPL.
- ▶ Once developed, a Vahana simulation can be **run parallelized** without additional actions from the developer.
- ▶ Models must be formulated as a **Graph Dynamic System** (GDS), which imposes certain constraints.

(Synchronous) Graph Dynamical Systems (GDS)

- ▶ A GDS involve a **finite graph** with a state assigned to each vertex and a **vertex function** that update the state of a vertex based on the states of its neighboring vertices.
- ▶ In a synchronous GDS (SyGDS), all vertex state updates happen **simultaneously**.
- ▶ A SyGDS can be interpreted as a **generalized cellular automata**, where the neighborhood is determined by the network instead of the position of the cell.
- ▶ When an ABM is implemented as a GDS, the **vertices represent the individual agents** of the model.

Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph** $G_t = (V_t, E_t)$.
- ▶ Vertices of **different types** $\delta_v \in \Delta_V$, and a mapping function $\pi_v : V_t \rightarrow \Delta_V$.
- ▶ Likewise edges of **different types** $\delta_e \in \Delta_E$, and a mapping function $\pi_e : E_t \rightarrow \Delta_E$.

Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph** $G_t = (V_t, E_t)$.
- ▶ Vertices of **different types** $\delta_v \in \Delta_V$, and a mapping function $\pi_v : V_t \rightarrow \Delta_V$.
- ▶ Likewise edges of **different types** $\delta_e \in \Delta_E$, and a mapping function $\pi_e : E_t \rightarrow \Delta_E$.
- ▶ For each δ_v and δ_e a different **state space** $\Theta_\delta = X_{\delta,1} \times \dots \times X_{\delta,n_\delta}$, where each $X_{\delta,i}$ is itself a set (like e.g. \mathbb{N}).
- ▶ At least one **transition (vertex) function** $f_{\delta_v,i} : G'_{v,t} \rightarrow (V'_{v,t+1}, E'_{v,t+1})$, where $G'_{v,t}$ is the 1-neighborhood of vertex $v \in G_t$.
- ▶ The source v_s and target v_t of an $e \in E'_{v,t+1}$ must be vertices of $G'_{v,t} \cup V'_{v,t+1}$.

Vahana's SyGDS Extension

We have ...

- ▶ A **directed graph** $G_t = (V_t, E_t)$.
- ▶ Vertices of **different types** $\delta_v \in \Delta_V$, and a mapping function $\pi_v : V_t \rightarrow \Delta_V$.
- ▶ Likewise edges of **different types** $\delta_e \in \Delta_E$, and a mapping function $\pi_e : E_t \rightarrow \Delta_E$.
- ▶ For each δ_v and δ_e a different **state space** $\Theta_\delta = X_{\delta,1} \times \dots \times X_{\delta,n_\delta}$, where each $X_{\delta,i}$ is itself a set (like e.g. \mathbb{N}).
- ▶ At least one **transition (vertex) function** $f_{\delta_v,i} : G'_{v,t} \rightarrow (V'_{v,t+1}, E'_{v,t+1})$, where $G'_{v,t}$ is the 1-neighborhood of vertex $v \in G_t$.
- ▶ The source v_s and target v_t of an $e \in E'_{v,t+1}$ must be vertices of $G'_{v,t} \cup V'_{v,t+1}$ **or a raster vertex**.

Constructing G_{t+1}

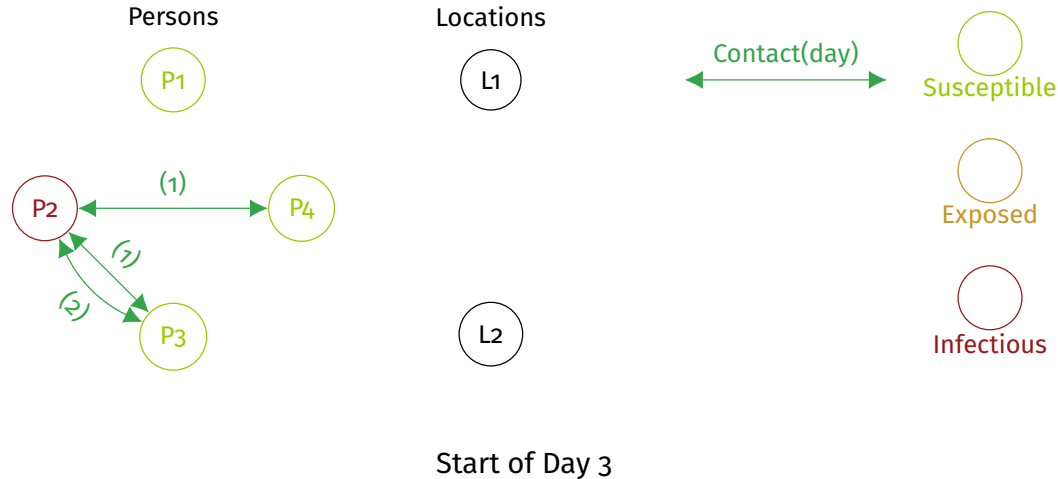
To construct G_{t+1} , we take the following actions:

- ▶ For each vertex $v \in V_t$ we **apply** $f_{\pi_v(v), i}$.
- ▶ **Combine** the results: $G_{t+1} = (\bigcup_{v \in V_t} V'_{v,t+1}, \bigcup_{v \in V_t} E'_{v,t+1})$.
- ▶ **Remove** all $e \in E_{t+1}$ with a source v_s or target v_t that is not in V_{t+1} .


It would be inefficient and cumbersome to require that even **invariant parts of the graph** must be reconstructed, so:

- ▶ Users specify which vertex/edge types are **mutable** during transitions.
- ▶ There exist an option to retain existing agents/edges of a type, **only adding new ones**.


Example: A epidemiological model



Example: A epidemiological model




Example: A epidemiological model



Locations: Infect persons & add contacts

Example: A epidemiological model



Persons: Update health status

REPL example

```
julia> show_agent(sim, Person; neighborstate = true)
Id / Local Nr: 0x01000000000085e35 / 548405
State:
  health=susceptible
  age=41
  quarantine=Quarantine(false, -1, -1)
Edge(s):
  MaybeEvent
    from:                edge.state:
    0x0200000000007691a MaybeEvent(0, ACTSTART, HOME)
    0x0200000000007691a MaybeEvent(32520, ACTEND, HOME)
    ... (16 not shown)
  Home
    from:
    0x0200000000007691a
  Knows
    from:                neighborstate:
    0x04000000000000001 capacity=0, traced=0, reported_infections=0,
```

Parallel Simulations

- ▶ Vahana.jl uses the **Message Passing Interface** (MPI) for communication across processes.
- ▶ But it **shields users** from complexities of MPI.
- ▶ At the end of the initialization phase, the actual graph is **automatically partitioned** and distributed to the processes.
- ▶ Each **process only knows a subset** of the graph representing simulation state.
- ▶ Vahana.jl uses **MPI-3 shared memory** to avoid copies of agent state on same node.

Hegselmann-Krause (HK) Opinion Model

- ▶ Finite number n of agents
- ▶ Time is **discrete**: $t = 0, 1, \dots$
- ▶ Real number $x_i(t)$ represents the **opinion** of agent i
- ▶ There is a **confidence bound** $\epsilon > 0$, opinions with a difference greater than ϵ are ignored.
- ▶ Opinions are updated **synchronously** according to

$$x_i(t+1) = \frac{1}{|\mathcal{N}_i(t)|} \sum_{j \in \mathcal{N}_i(t)} x_j(t)$$

$$\text{where } \mathcal{N}_i(t) = \{j : \|x_j(t) - x_i(t)\| \leq \epsilon\}$$